



# Distributed Triangle Counting in the Graphulo Matrix Math Library

Dylan Hutchison, University of Washington, dhutchis@cs.washington.edu




## An open source library to orchestrate server-side graph processing in the Apache Accumulo database

- Graphulo uses server-side iterators (a la BigTable) to compute the GraphBLAS kernels on database tables:
  - Matrix Multiply
  - Matrix Reduce
  - Apply function
  - Matrix Extraction (Row/Column filtering)
  - Element-wise Multiply
  - Element-wise Add
  - Matrix Assignment
- Advantages of in-database computation include
  - Data Locality (Graphulo brings compute to data)
  - Infrastructure Reuse (no sharing hardware between systems)
  - Accumulo Features (Accumulo as a Big Index, Distributed Execution)



## Static Graph Challenge

- Finding the k-Truss subgraph
  - Implemented in HPEC '16
- Triangle Counting
  - This poster



## Two Triangle Counting Algorithms

**Input:** Upper triangle of unweighted adjacency matrix  $A$   
**Output:** Number of triangles  $t$

```
1  $T = A$  // clone  $A$  to  $T$ 
2  $T = T + \text{triu}(A^T A)$  // upper triangle of matrix multiply
  // custom multiply:  $a \otimes b = 2$  if  $a = b = 1$ , otherwise 0
3  $T(T \% 2 == 0) = 0$  // filter to odd entries
4  $t = \text{sum}((T - 1)/2)$ 
```

**Algorithm 2:** Graphulo Adjacency-only triangle counting

**Input:** Lower triangle of unweighted adjacency matrix  $A$   
**Input:** Unweighted incidence matrix  $E$   
**Output:** Number of triangles  $t$


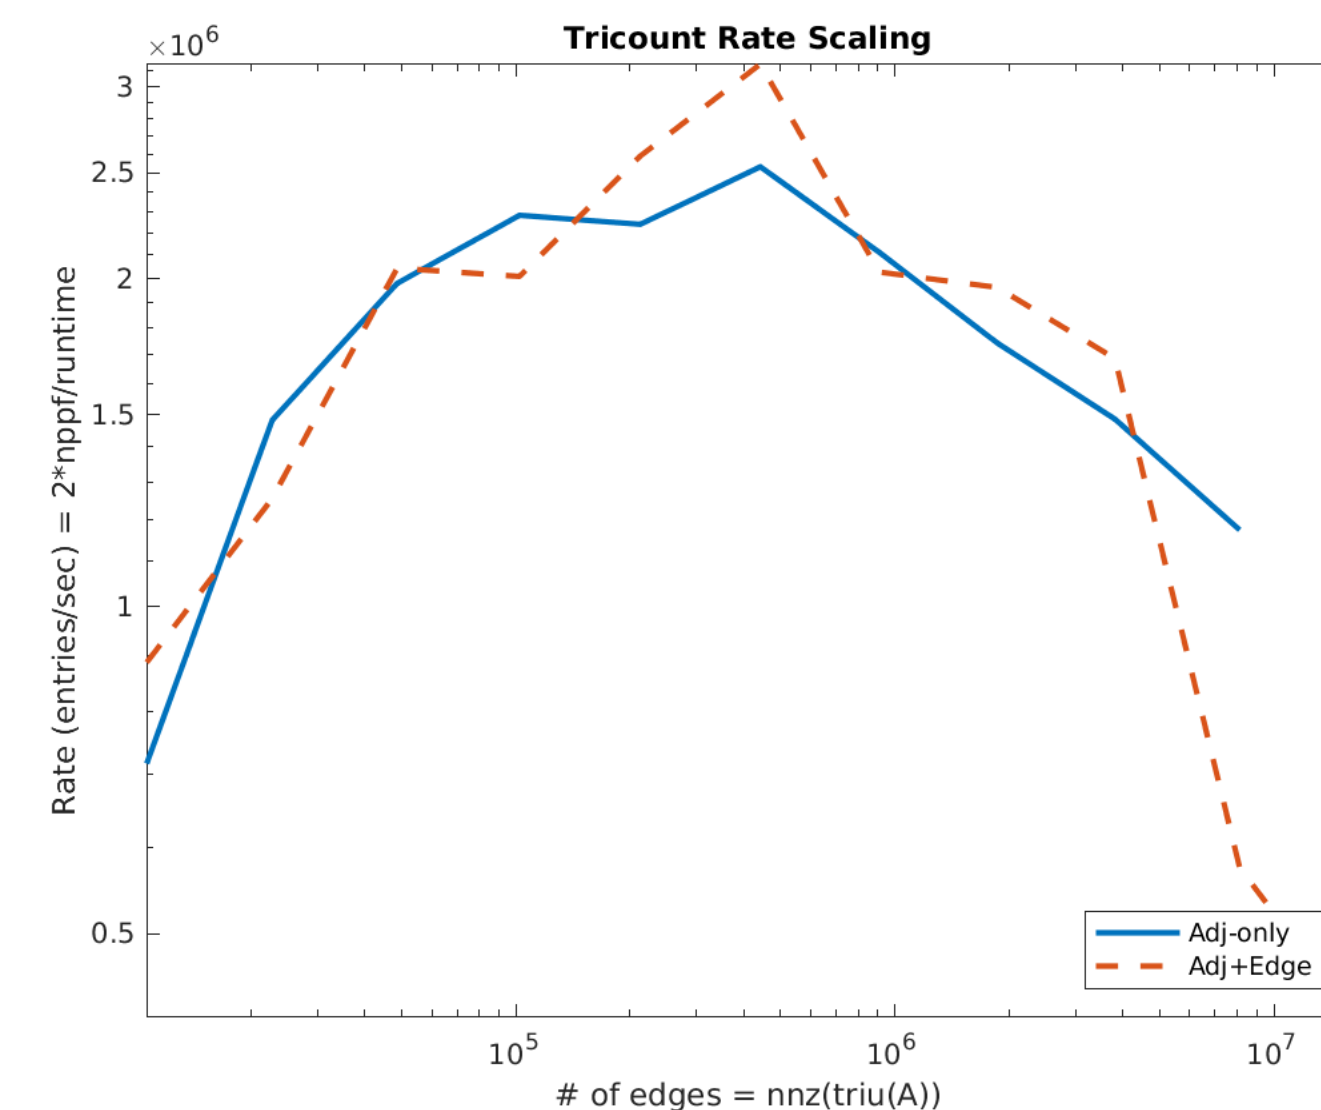
```
1  $T = \text{triu}(A^T E)$  // upper triangle of matrix multiply
2  $t = \text{sum}(T == 2)$  // count the entries of  $T$  equal to 2
```

**Algorithm 3:** Graphulo Adj.+Incidence triangle counting

## Triangle Counting Experiment on Power Law Graphs

Question: How do the two triangle counting implementations compare in performance?

Baseline: Single-node in-memory MATLAB implementation



## Discussion

- MATLAB Baseline exceeds memory after scale 15 (520k edges); Baseline performs faster while it fits in memory
- Similar performance for both Graphulo algorithms
- Phase transition between scale 15 and 16 (520k and 1.05M edges)
  - Scale  $\leq 15$  – Reduction step is bottleneck
  - Scale  $\geq 16$  – Matrix multiply step is bottleneck
  - Peak processing rate achieved at the phase transition point
- Data skew (from power law graphs) limits performance
  - Poor load balancing; one machine receives a disproportionately large workload
  - Permuting rows and columns helps; however, some machine must receive the “super-node”—the highest degree row
- Proposal: A hybrid algorithm to address skew
  - Handle high-degree rows separately with an inner-product matrix multiply
  - Handle low-degree rows with an outer-product matrix multiply

## Experiment Details

- m3.xlarge Amazon nodes, each 15 GB mem, 4 vCPU, 2x40 GB SSD
- 8 workers, 3 coordinators, 1 monitor
- Graph500 power law matrix generator,  $2^{10}$  to  $2^{20}$  rows, 16 nonzeros/row

## Past Work

- Showed Graphulo faster than single-node in-memory LA packages on MxM (HPEC '15)
- Confirmed results for more complex I/O-bound, single-pass graph analytics (IPDPS '15, HPEC '16)
- Verified Graphulo scales with Accumulo as cluster size increases (HPEC '16)
- Implemented GraphBLAS, Jaccard, k-Truss; showed Graphulo is best for I/O-bound, single-pass analytics (HPEC '16)
- Built LaraDB atop Graphulo; compared to MapReduce on MxM (BeyondMR '17)

## Future Work

- Measure how well a hybrid algorithm addresses data skew
- Explore the “export to an external system” strategy, as used in polystores such as Myria and BigDAWG
- Benchmark on real-world data, which likely has less skew
- Address programmability barriers to graph algorithms via code generation

