



# A Minimalist Kernel for Linear and Relational Algebra

---

Dylan Hutchison, Bill Howe, Dan Suciu  
BeyondMR @SIGMOD, 19 May 2017

Contact: [dhutchis@cs.washington.edu](mailto:dhutchis@cs.washington.edu)

UNIVERSITY *of* WASHINGTON

W

# Related Work

2017  
LARADB

## LaraDB: A Minimalist Kernel for Linear and Relational Algebra Computation

Dylan Hutchison, Bill Howe, Dan Suciu

- Sensor Query on *Array of Things* data
- Logical + Physical Algebra
- Properties / Optimizations
- LaraDB Implementation
- Experiments

LARA: A Key-Value Algebra underlying Arrays and Relations

Dylan Hutchison, Bill Howe, Dan Suciu

University of Washington

Draft: April 1, 2016

2016

Abstract

- Logical Algebra
- Properties / Optimizations
- Translations

Data  
when  
systems  
disparities  
algebraic  
operations  
equivalence  
role of

Many  
other  
able  
seen  
ifies  
the  
and  
the

ABSTRACT  
Advanced  
processing  
processing  
ysis tech  
ressed  
while m  
iteration  
analysis  
tems. T  
tivity but  
sharing of physical data layouts (e.g., partitioning) and data

- Vision paper
- Hypothetical Sensor Query
- Properties / Optimizations

## Scalable Linear Algebra on a Relational Database System

Shangyu Luo<sup>\*</sup>, Zekai J. Gao<sup>\*</sup>, Michael Gutinov<sup>†</sup>, Luis L. Perez<sup>‡</sup>, Christopher Jermaine<sup>\*</sup>

<sup>\*</sup>Rice University, {s445, jacobgao, cm4}@rice.edu, <sup>†</sup>Rice University, lperez@gmail.com

<sup>‡</sup>University of Texas at Austin, cjermaine@cs.utexas.edu

## The BUDS Language for Distributed Bayesian Machine Learning

Zekai J. Gao, Shangyu Luo, Luis L. Perez, Chris Jermaine  
Rice University

## Towards Linear Algebra over Normalized Data

Lingjiao Chen<sup>1</sup>, Arun Kumar<sup>2</sup>

<sup>1</sup>University of Wisconsin-Madison

Jeffrey Naughton<sup>3</sup>

<sup>3</sup>University of California, San Diego

Jignesh M. Patel<sup>4</sup>  
<sup>4</sup>Google

2017

Shangyu Luo<sup>\*</sup>, Zekai J. Gao<sup>\*</sup>, Michael Gutinov<sup>†</sup>, Luis L. Perez<sup>‡</sup>, Christopher Jermaine<sup>\*</sup>

<sup>\*</sup>Rice University, {s445, jacobgao, cm4}@rice.edu, <sup>†</sup>Rice University, lperez@gmail.com

<sup>‡</sup>University of Texas at Austin, cjermaine@cs.utexas.edu

## The BUDS Language for Distributed Bayesian Machine Learning

Zekai J. Gao, Shangyu Luo, Luis L. Perez, Chris Jermaine  
Rice University

## Towards Linear Algebra over Normalized Data

Lingjiao Chen<sup>1</sup>, Arun Kumar<sup>2</sup>

<sup>1</sup>University of Wisconsin-Madison

Jeffrey Naughton<sup>3</sup>

<sup>3</sup>University of California, San Diego

Jignesh M. Patel<sup>4</sup>  
<sup>4</sup>Google

- Extend SQL with LA types, UDFs on *SimSQL*
- Compile MATLAB-like LA to SQL in *BUDS*
- *Morpheus*: Push LA into matrix multiply for ML

## Bridging the Gap: Towards Optimization Across Linear and Relational Algebra

Andreas Kunft

Alexander Alexandrov

Asterios Katsifodimos

Volker Markl

TU Berlin  
first.last@tu-berlin.de

Shangyu Luo<sup>\*</sup>, Zekai J. Gao<sup>\*</sup>, Michael Gutinov<sup>†</sup>, Luis L. Perez<sup>‡</sup>, Christopher Jermaine<sup>\*</sup>

<sup>\*</sup>Rice University, {s445, jacobgao, cm4}@rice.edu, <sup>†</sup>Rice University, lperez@gmail.com

<sup>‡</sup>University of Texas at Austin, cjermaine@cs.utexas.edu

## The BUDS Language for Distributed Bayesian Machine Learning

Zekai J. Gao, Shangyu Luo, Luis L. Perez, Chris Jermaine  
Rice University

## Towards Linear Algebra over Normalized Data

Lingjiao Chen<sup>1</sup>, Arun Kumar<sup>2</sup>

<sup>1</sup>University of Wisconsin-Madison

Jeffrey Naughton<sup>3</sup>

<sup>3</sup>University of California, San Diego

Jignesh M. Patel<sup>4</sup>  
<sup>4</sup>Google

- Extend SQL with LA types, UDFs on *SimSQL*
- Compile MATLAB-like LA to SQL in *BUDS*
- *Morpheus*: Push LA into matrix multiply for ML

## Bridging the Gap: Towards Optimization Across Linear and Relational Algebra

Andreas Kunft

Alexander Alexandrov

Asterios Katsifodimos

Volker Markl

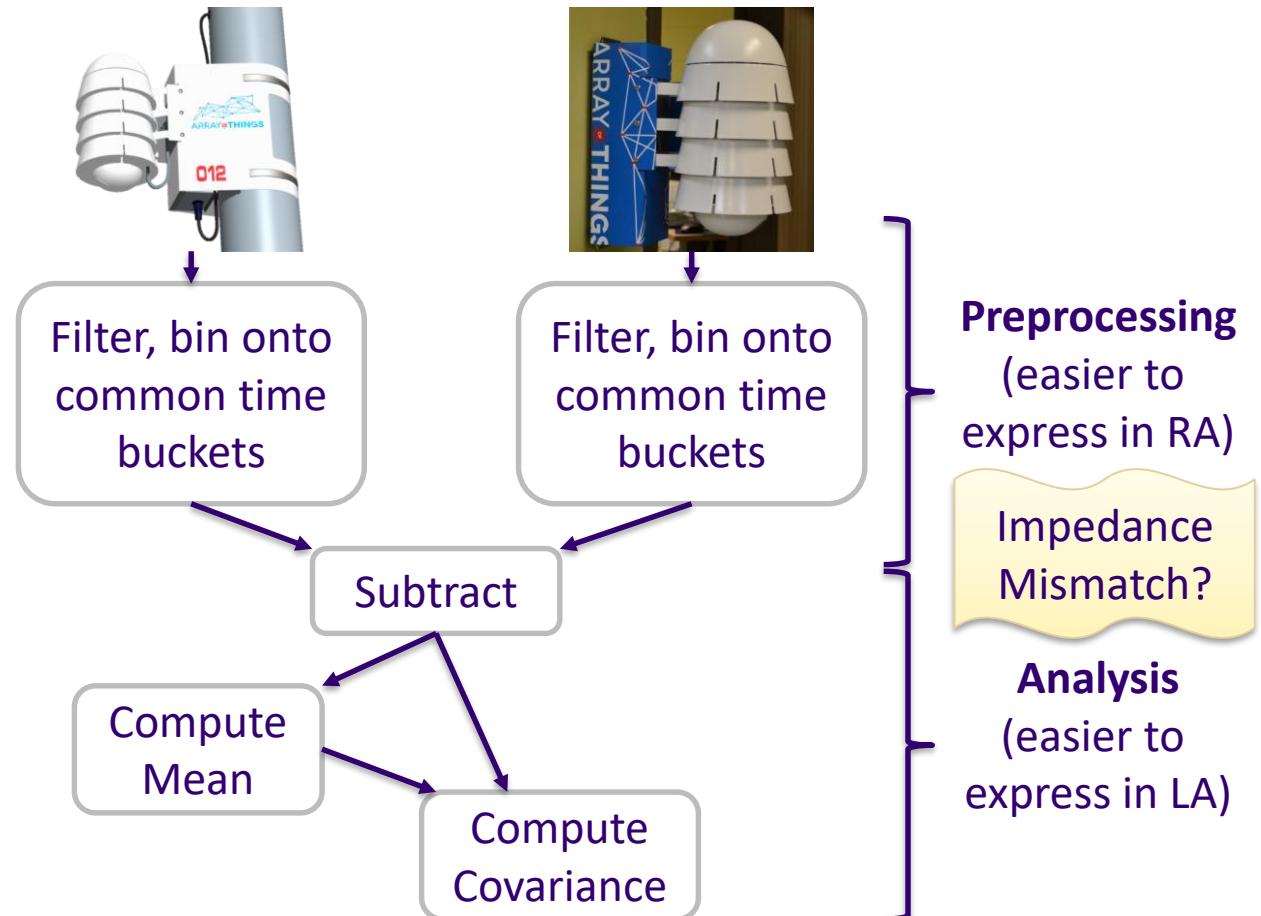
TU Berlin  
first.last@tu-berlin.de

# Motivating Use Case: Mean & Covariance of sensor differences

Array of Things  
Sensor Data

t	c	v
466	temp	55.2
466	hum	40.1
492	temp	56.3
492	hum	35.0
528	temp	56.5

Collected in CSV files



# Bin query: easy in RA, awkward in LA

t	c	v
466	temp	55.2
466	hum	40.1
492	temp	<b>56.3</b>
492	hum	35.0
528	temp	<b>56.5</b>



**RA**  
SELECT bin(t) AS t', c, avg(v) AS v  
GROUP BY t', c

t'	c	v
460	temp	55.2
460	hum	40.1
520	temp	<b>56.4</b>
520	hum	35.0

**LA**

Multiply:

$$\begin{matrix} 466 & 492 & 528 \\ 460 & [1 & 1 & 1] \end{matrix}$$



$$\begin{matrix} temp & hum \\ 466 & [55.2 & 40.1] \\ 492 & [56.3 & 35.0] \\ 528 & [56.5 & 35.0] \end{matrix} = \begin{matrix} temp & hum \\ 460 & [55.2 & 40.1] \\ 520 & [56.4 & 35.0] \end{matrix}$$

using *avg* on added elements

# Covariance query: easy in LA, awkward in RA

$$X = \begin{bmatrix} x_{11} & x_{12} \\ x_{21} & x_{22} \\ x_{31} & x_{32} \end{bmatrix}$$

## LA

```
N = size(X, 1);  
M = mean(X, 1);  
C = X'*X / N - M'*M;
```

## RA

(Generated SQL statements for each entry)

```
T = SELECT FROM X sum(1.0) AS N,  
sum(X1) AS M1, sum(X2) AS M2, ..., sum(Xd) AS Md,  
sum(X1*X1) AS Q11, sum(X1*X2) AS Q12, ...,  
sum(Xd-1*Xd) AS Q(d-1)d, sum(Xd*Xd) AS Qdd  
C = SELECT FROM T  
(1 AS i, 1 AS j, Q11/N - M1*M1 AS v) UNION  
(1 AS i, 2 AS j, Q12/N - M1*M2 AS v) UNION  
...
```

# Properties, easier in LA or RA

---

- >  $(A^T A)^T = A^T A$  [inner product is symmetric]
  - RA version is verbose (renames, joins, aggregates)
  - Impact: eliminate half of computation
  
- >  $A \bowtie_i B = \sigma_{i=i} (A \times B)$  [cross product to join]
  - LA version is verbose (tensor products and contractions)
  - Impact: eliminate large intermediates

W

# Is a common algebra possible?

---

Evidence:

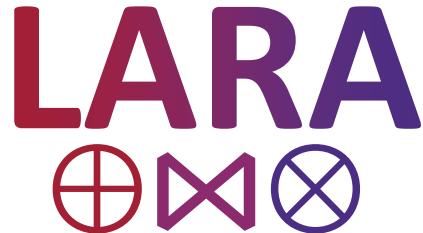
- > Overlapping expressiveness
  - Matrix Multiply = GroupBy o Apply o Join
- > Emergence of hybrid systems
  - Spark, Myria, MADlib, SystemML, Weld, ...
  - Systems with RA and LA interfaces or data models



# Roadmap

---

- > Motivation: Sensor query
  - Sometimes RA is better, sometimes LA
- **LARA Algebra**
  - 3 Operators & Properties
- > LARADB on *Sorted Distributed Maps*
  - Physical operators & optimizations
- > Experiments
  - Optimizations: LARA affords impactful optimizations
  - Performance: LARA affords competitive performance



## Objects: *Associative Tables*

Total functions from keys to values with finite support

		Attributes	
		Keys	Values
		[0]	[“”]
$k_1$	$k_2$	$v_1$	$v_2$
a	37	7	'dan'
a	20	0	"
b	25	0	'dylan'
b	20	2	'bill'

Annotations:

- Brackets labeled "Default Values" point to the header cells [0] and [“”].
- A bracket labeled "Each row is a tuple" points to the first four data rows.
- A bracket labeled "Support" points to the last two data rows.

## Operators:

Parameterized by  
UDFs:  $\otimes$ ,  $\oplus$ , f

*Join*



“join keys,  
multiply values”

*Union*



“group by keys,  
add values”

*Extension*

$\text{ext}_f$

“flatmap”

Join and Union adapted from:  
M. Spight and V. Tropashko.  
*First steps in relational lattice*. 2006.

Ext is similar to  
monadic bind,  
without deleting keys

# Join with $\otimes$ :

## Join keys, multiply values

		[0]	[0]		
a	c	x	z		
<hr/>					
a <sub>1</sub>	c <sub>1</sub>	11	1		
a <sub>1</sub>	c <sub>2</sub>	12	2		
a <sub>2</sub>	c <sub>1</sub>	13	3		
a <sub>3</sub>	c <sub>3</sub>	14	4		

 $\bowtie_{\otimes}$ 

		[0]	[0]		
c	b	z	y		
<hr/>					
c <sub>1</sub>	b <sub>1</sub>	5	15		
c <sub>2</sub>	b <sub>1</sub>	6	16		
c <sub>2</sub>	b <sub>2</sub>	7	17		
c <sub>4</sub>	b <sub>1</sub>	8	18		

=

			[0 $\otimes$ 0 = 0]	
a	c	b	z	
<hr/>				
a <sub>1</sub>	c <sub>1</sub>	b <sub>1</sub>	1 $\otimes$ 5	
a <sub>1</sub>	c <sub>2</sub>	b <sub>1</sub>	3 $\otimes$ 5	
a <sub>1</sub>	c <sub>2</sub>	b <sub>2</sub>	2 $\otimes$ 6	
(a <sub>3</sub>	c <sub>3</sub>	b <sub>1</sub>	4 $\otimes$ 0 = 0)	

$$A \bowtie_{\otimes} B : a, c, b \rightarrow z : 0^z \otimes 0^z$$

$$(A \bowtie_{\otimes} B)(a, c, b) := [z : \pi_z A(a, c) \otimes \pi_z B(c, b)]$$

Requires (see paper for general case):

$$v_A \otimes 0 = 0 \otimes v_B = 0 \otimes 0 = 0$$

# Union with $\oplus$ : Group by keys, sum values

a	c	[0]	[0]	$\sum_{\oplus}$	c	b	[0]	[0]	=	c	[0]	[0 $\oplus$ 0 = 0]	[0]
a <sub>1</sub>	c <sub>1</sub>	11	1		c <sub>1</sub>	b <sub>1</sub>	5	15		c <sub>1</sub>	x	11 $\oplus$ 13	15
a <sub>1</sub>	c <sub>2</sub>	12	2		c <sub>2</sub>	b <sub>1</sub>	6	16		c <sub>2</sub>	z	1 $\oplus$ 3 $\oplus$ 5	
a <sub>2</sub>	c <sub>1</sub>	13	3		c <sub>2</sub>	b <sub>2</sub>	7	17		c <sub>3</sub>	y	16 $\oplus$ 17	
a <sub>3</sub>	c <sub>3</sub>	14	4		c <sub>4</sub>	b <sub>1</sub>	8	18		c <sub>4</sub>		4	0
												8	18

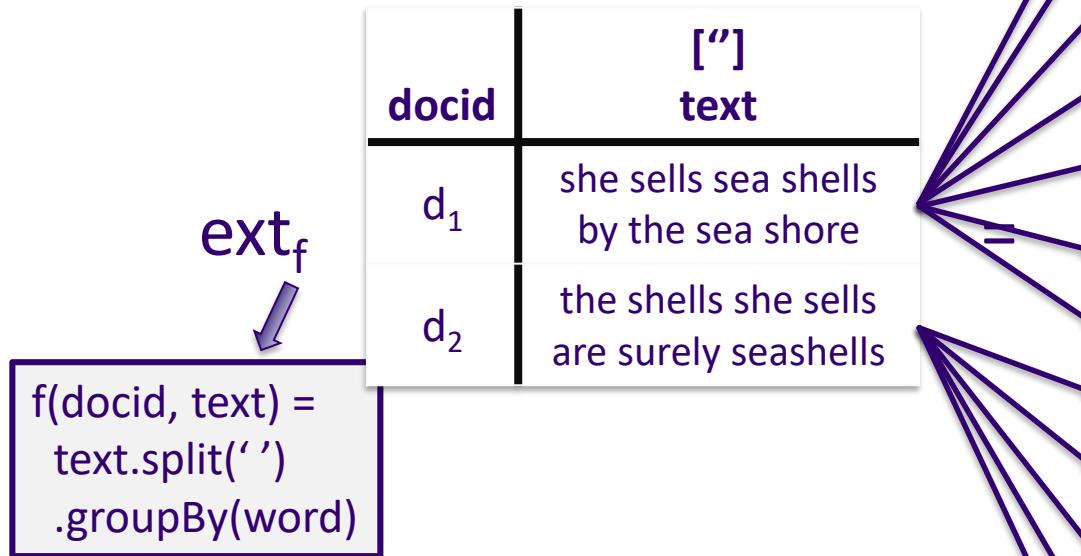
$$A \sum_{\oplus} B : c \rightarrow x, z, y : 0^x, 0^z, 0^y$$

$$(A \sum_{\oplus} B)(c) := [x: \bigoplus_a \pi_x A(a, c),$$

$$z: \bigoplus_a \pi_z A(a, c) \oplus \bigoplus_b \pi_z B(c, b), y: \bigoplus_b \pi_y B(c, b)]$$

Requires:  
 $v \oplus 0 = 0 \oplus v = v$

# Ext with f: Flatmap



Produces new key 'word'

docid	word	[0] value
$d_1$	she	1
$d_1$	sells	1
$d_1$	sea	1
$d_1$	shells	1
$d_1$	by	1
$d_1$	the	1
$d_1$	shore	1
$d_2$	the	1
$d_2$	shells	1
$d_2$	she	1
$d_2$	sells	1
$d_2$	are	1
...	...	...

$$\text{ext}_f A : a, c, k' \rightarrow v' : 0'$$

$$(\text{ext}_f A)(a, c, k') := f(a, c, A(a, c))(k')$$

Requires that  $f$  behave same  
on default value input  
UNIVERSITY of WASHINGTON

# Summary: Union, Join, Ext

	Key Types	Value Types	Support	
<b>Union</b> ( $A \boxtimes_{\oplus} B$ )	$= K_A \cap K_B$	$= V_A \cup V_B$	$\subseteq S_A \cup S_B$	<i>Duality</i>
<b>Join</b> ( $A \bowtie_{\otimes} B$ )	$= K_A \cup K_B$	$= V_A \cap V_B$	$\subseteq S_A \cap S_B$	
<b>Ext</b> ( $\text{ext}_f A$ )	extended by f	set by f	$\subseteq S_A \times S_f$	

For Support, ' $\subseteq$ ' becomes '=' if  $\oplus$  is zero-sum-free or  
 $\otimes$  has zero-product-property (similar for f)

# LARA Properties

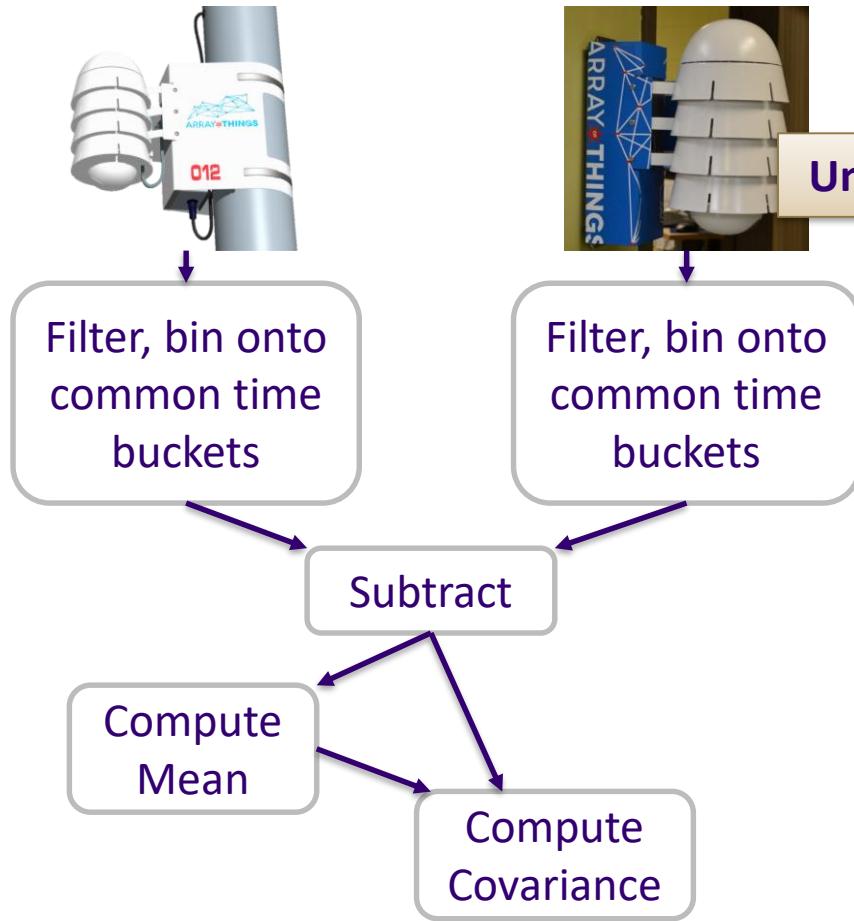
---

- If  $\oplus$  or  $\otimes$  is associative, commutative, or idempotent,
  - then so is Union or Join
- If  $\otimes$  distributes over  $\oplus$ ,
  - Push Union into Join
  - Distribute Join over Union

➤ See paper for more properties:

$$\begin{aligned} & \text{tr}(ABC) && \text{tr defn.} \\ &= \boxtimes_+ \text{ext}_{i=l}(ABC) && \text{tr defn.} \\ &= \boxtimes_+ \boxtimes_{i=l}^{\text{i}, \text{l}} (\boxtimes_+^{\text{i}, \text{k}}(A_{ij} \bowtie_{\otimes} B_{jk}) \bowtie_{\otimes} C_{kl}) && ABC \text{ defn.} \\ &= \boxtimes_+ \boxtimes_{i=l}^{\text{i}, \text{l}} \text{ext}_{i=l}(\boxtimes_+^{\text{i}, \text{k}}(A_{ij} \bowtie_{\otimes} B_{jk}) \bowtie_{\otimes} C_{kl}) && \text{push ext into } \boxtimes \\ &= \boxtimes_+ \text{ext}_{i=l}(\boxtimes_+^{\text{i}, \text{k}}(A_{ij} \bowtie_{\otimes} B_{jk}) \bowtie_{\otimes} C_{kl}) && \text{combine } \boxtimes \\ &= \boxtimes_+(\boxtimes_+^{\text{i}, \text{k}}(A_{ij} \bowtie_{\otimes} B_{jk}) \bowtie_{\otimes} C_{ki}) && \text{apply ext} \\ &= \boxtimes_+ \boxtimes_+^{\text{i}, \text{k}}(A_{ij} \bowtie_{\otimes} B_{jk} \bowtie_{\otimes} C_{ki}) && \text{distr. } \bowtie \text{ into } \boxtimes \\ &= \boxtimes_+(A_{ij} \bowtie_{\otimes} B_{jk} \bowtie_{\otimes} C_{ki}) && \text{combine } \boxtimes \\ &= \boxtimes_+(B_{jk} \bowtie_{\otimes} C_{ki} \bowtie_{\otimes} A_{ij}) && \text{commute } \bowtie_{\otimes} \\ &= \dots // \text{reversing the above steps} \\ &= \text{tr}(BCA) \end{aligned}$$

# Sensor Query: SQL/MATLAB to Lara



**Ext**

```
A, B = LOAD('s1.csv'), LOAD('s2.csv')  
A' = SELECT bin(t) AS t', c, avg(v) AS v  
FROM A WHERE 460 <= t <= 860  
GROUP BY t', c
```

**Union**

```
B' = ... // repeat above
```

**Ext**

```
// change perspective from RA to LA
```

**Join**

```
X = A - B;
```

**Ext, Union**

```
N = size(X, 1);
```

```
M = mean(X, 1);
```

```
disp(M);
```

**Join**

```
U = X - repmat(M, N, 1);
```

```
C = U.' * U ./ (N - 1);
```

**Rename (Ext), Join, Union**

```
disp(C);
```



# Roadmap

---

- > Motivation: Sensor query
  - Sometimes RA is better, sometimes LA
- > LARA Algebra
  - 3 Operators & Properties
- **LARADB on *Sorted Distributed Maps***
  - **Physical operators & optimizations**
- > Experiments
  - Optimizations: LARA affords impactful optimizations
  - Performance: LARA affords competitive performance

# Implementing LARA on relevant systems with Sorted Partitioned Maps

An abstraction for:

- > Relational systems
  - with primary key
- > NoSQL key-value systems
- > Array systems

k	v
$k_1$	v
$k_2$	v
$k_3$	v
$k_4$	v
$k_5$	v
$k_6$	v
$k_7$	v

Partition

Partition

**Map:** O(1) access by key

**Sorted:** O(1) access to “next” element

**Partitioned:** Scale-out

# Physical LARA on Sorted Partitioned Maps

*Sorted Associative Table:*

Ordered key attributes  $[t, c]$

Physical Operators:

- MergeJoin
- MergeUnion
- Ext
- **Sort**
- (Load)
- (Store)

Physical Optimizations



(BigTable)

	t	c	[ $\perp$ ] v
Partition	440	temp	54.0
	440	hum	40.8
	466	temp	55.2
	466	hum	40.1
Partition	492	temp	56.3
	492	hum	35.0
	528	temp	56.5

Accumulo Row  
Column Qualifier  
Value

LARADB prototype in Accumulo  
on Graphulo iterators



# Roadmap

---

- > Motivation: Sensor query
  - Sometimes RA is better, sometimes LA
- > LARA Algebra
  - 3 Operators & Properties
- > LARADB on *Sorted Distributed Maps*
  - Physical operators & optimizations

## → Experiments

- Optimizations: LARA affords impactful optimizations
- Performance: LARA affords competitive performance

# Optimization Experiment

---

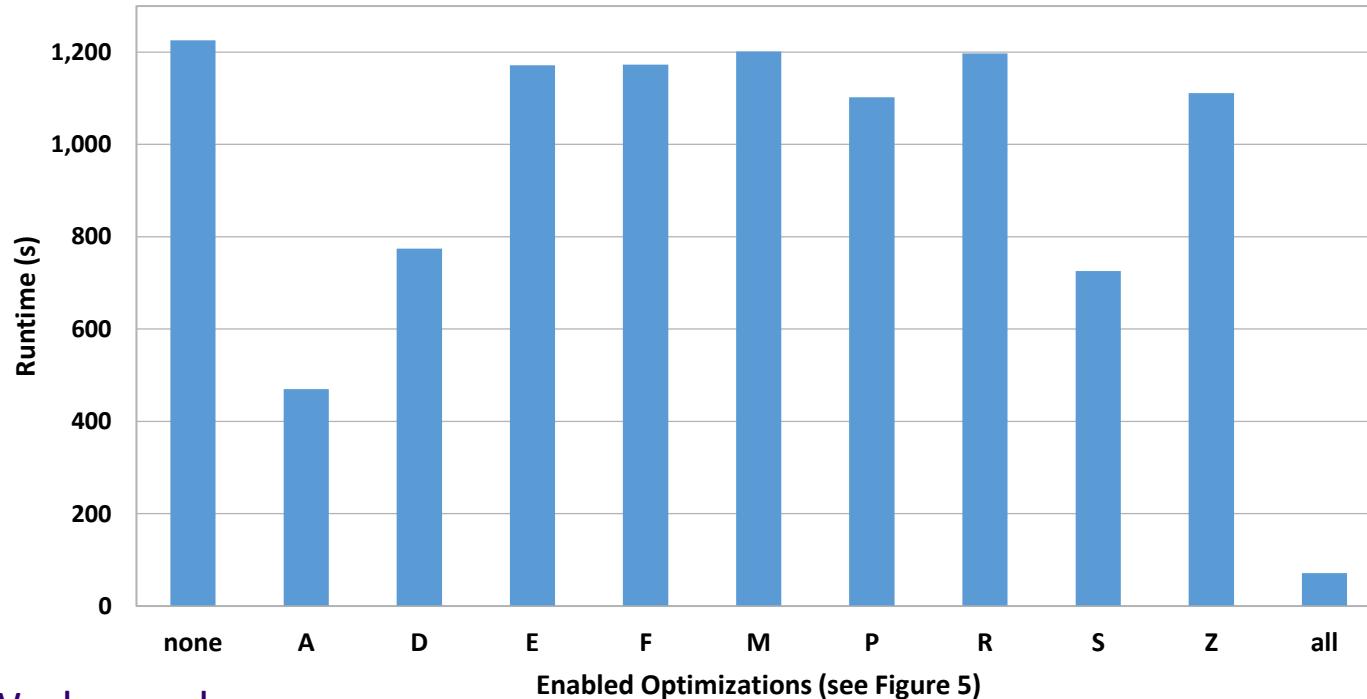
*Does LARA afford optimizations  
that impact performance?*

Approach: Benchmark Sensor query on LARADB.  
Implement optimizations. Plot runtime reduction.

Cluster: Amazon EC2 m3.large  
(2 vCores, 7.5 GB mem, SSD per node)

# Optimization Experiment

Effect of Optimizations on Sensor Calculation Runtime



- 4 Worker nodes
- 1.5 months' sensor data
  - Partitioned in 3 day segments
  - 1.2 GB raw; 60 MB after filter to 1 month, project, compress in HDFS

# Performance Experiment

---

*Does LARA afford  
competitive performance  
when implemented?*

Approach: Benchmark Matrix Multiply on data in Accumulo

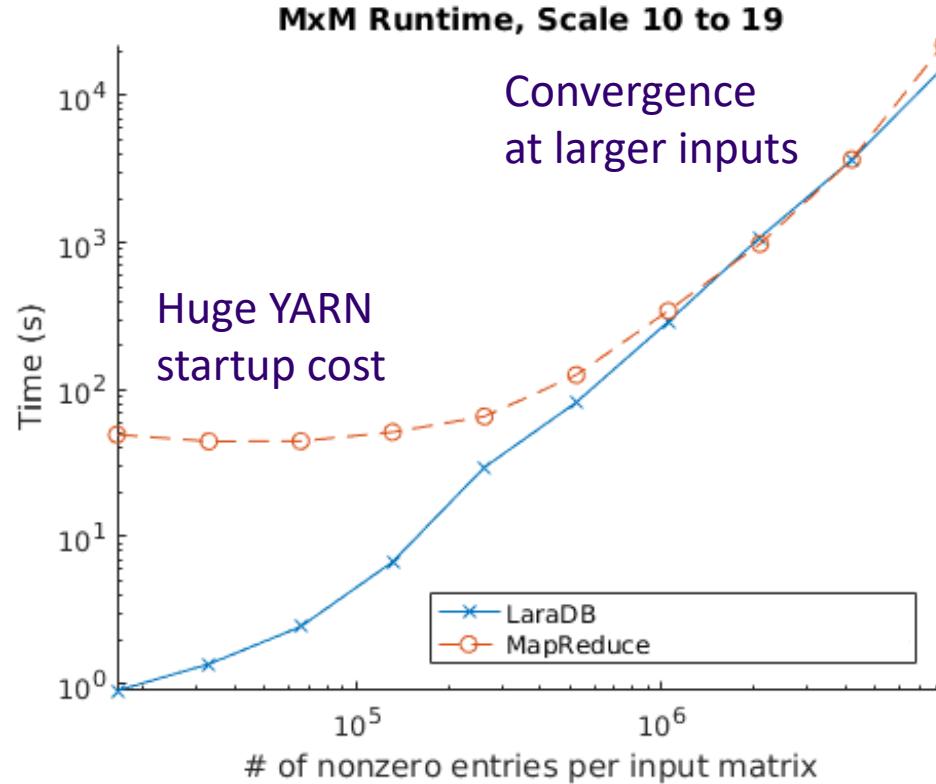
Engines: LARADB and Hadoop MapReduce

Task: Matrix-Matrix Multiply (Graph500 power law matrices)

- MxM is a building block of both LA and RA (join, aggregate)

Cluster: Amazon EC2 m3.large  
(2 vCores, 7.5 GB mem, SSD per node)

# Performance Experiment

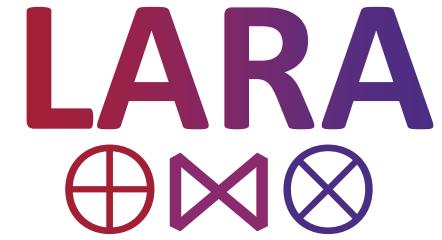


- 8 Worker nodes
- $16 \times 2^{\text{SCALE}}$  entries, generated from Graph500, partitioned evenly on input rows
- Scale 19 MxM forms  $8 \times 10^9$  partial products, skewed due to power law

# Recap

---

- > Motivation: Sensor query
  - Sometimes RA is better, sometimes LA
- > LARA Algebra
  - 3 Operators & Properties
- > LARADB on *Sorted Distributed Maps*
  - Physical operators & optimizations
- > Experiments
  - Optimizations: LARA affords impactful optimizations
  - Performance: LARA affords competitive performance



*Join*      *Union*

$\bowtie_{\otimes}$        $\Sigma_{\oplus}$

*Extension*

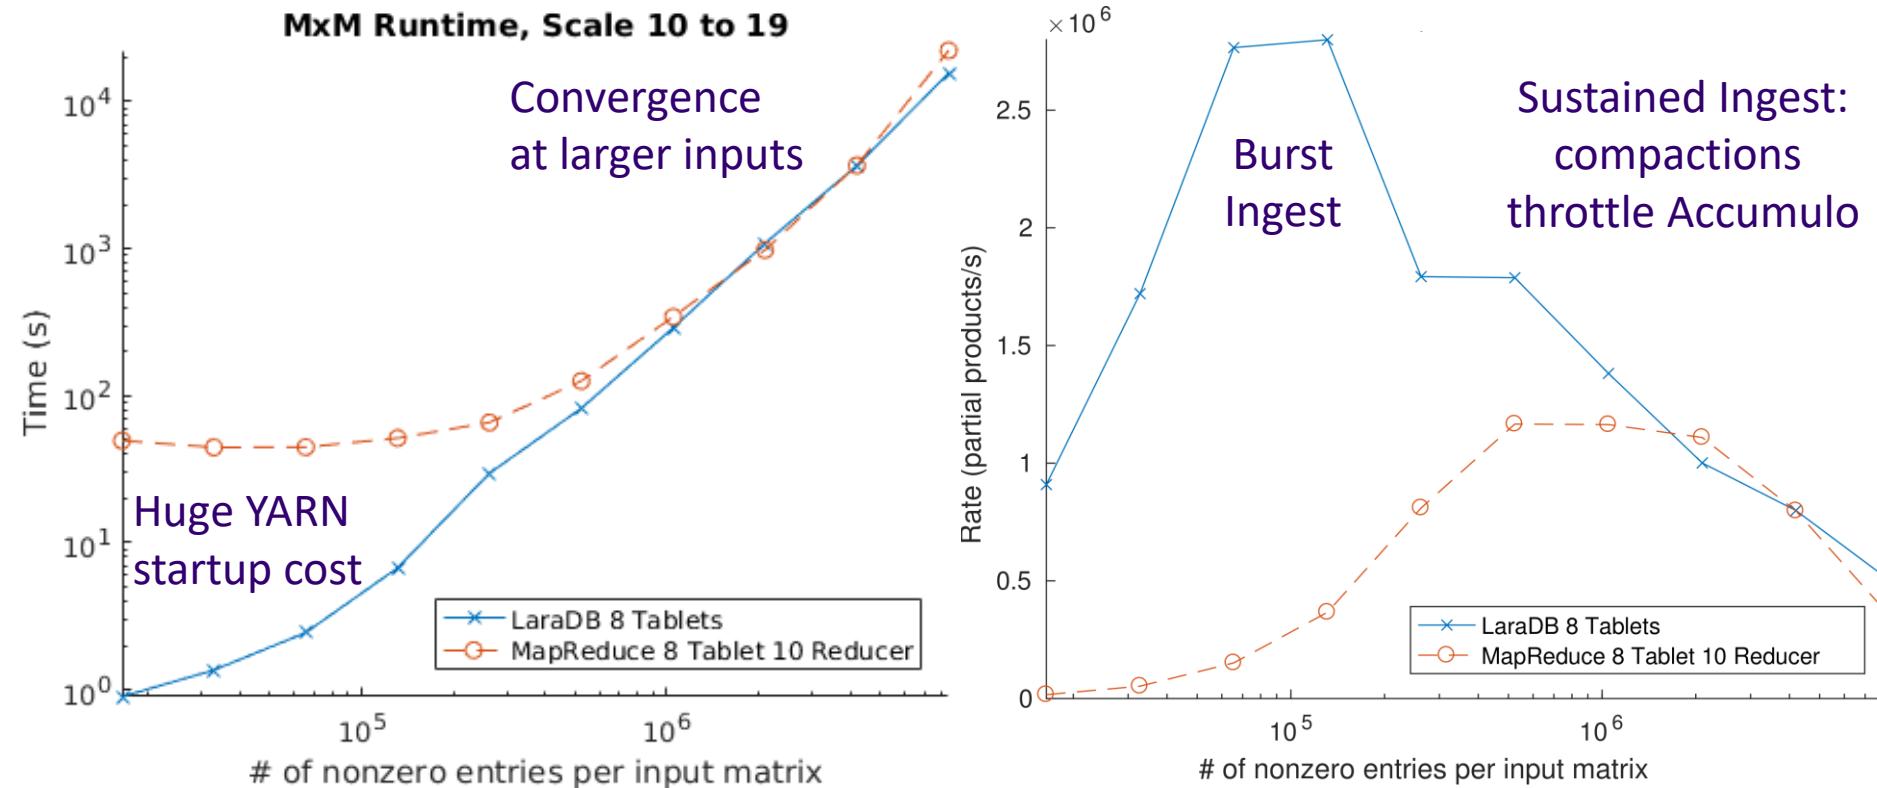
$\text{ext}_f$

Additional thanks to  
David Maier, Tim Mattson,  
Jeremy Kepner, Vijay Gadepally



# Backup Slides

# Competitiveness Experiment



- 8 Worker nodes
- $16 \times 2^{\text{SCALE}}$  entries, generated from Graph500, partitioned evenly on input rows
- Scale 19 MxM forms  $8 \times 10^9$  partial products, skewed due to power law

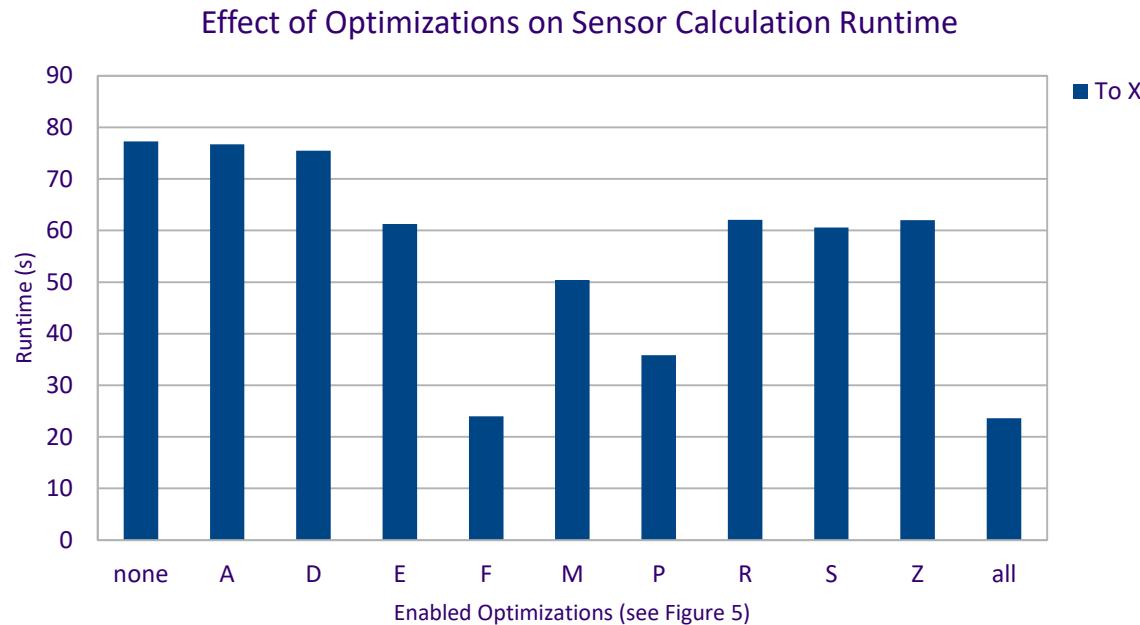
# Data for Competitiveness Experiment

Scale	Entries in Table C		LARADB		MapReduce 10 Reducers	
	PartialProducts	AfterSum	Time (s)	Rate (pp/s)	Time (s)	Rate (pp/s)
10	$8.16 \times 10^5$	$2.69 \times 10^5$	$8.98 \times 10^{-1}$	$9.09 \times 10^5$	$4.95 \times 10^1$	$1.65 \times 10^4$
11	$2.34 \times 10^6$	$7.90 \times 10^5$	1.36	$1.72 \times 10^6$	$4.45 \times 10^1$	$5.26 \times 10^4$
12	$6.77 \times 10^6$	$2.41 \times 10^6$	2.45	$2.77 \times 10^6$	$4.47 \times 10^1$	$1.52 \times 10^5$
13	$1.89 \times 10^7$	$7.04 \times 10^6$	6.74	$2.80 \times 10^6$	$5.15 \times 10^1$	$3.66 \times 10^5$
14	$5.32 \times 10^7$	$2.34 \times 10^7$	$2.97 \times 10^1$	$1.79 \times 10^6$	$6.57 \times 10^1$	$8.10 \times 10^5$
15	$1.47 \times 10^8$	$6.38 \times 10^7$	$8.21 \times 10^1$	$1.79 \times 10^6$	$1.26 \times 10^2$	$1.17 \times 10^6$
16	$4.00 \times 10^8$	$2.30 \times 10^8$	$2.89 \times 10^2$	$1.38 \times 10^6$	$3.44 \times 10^2$	$1.16 \times 10^6$
17	$1.09 \times 10^9$	$9.19 \times 10^8$	$1.09 \times 10^3$	$1.00 \times 10^6$	$9.80 \times 10^2$	$1.11 \times 10^6$
18	$2.94 \times 10^9$	$2.55 \times 10^9$	$3.67 \times 10^3$	$7.99 \times 10^5$	$3.68 \times 10^3$	$7.97 \times 10^5$
19	$7.84 \times 10^9$	$6.58 \times 10^9$	$1.56 \times 10^4$	$5.02 \times 10^5$	$2.23 \times 10^4$	$3.52 \times 10^5$

**Table 1: 8-worker  $A^\top B$  experiment Runtimes, Rates, and Statistics**

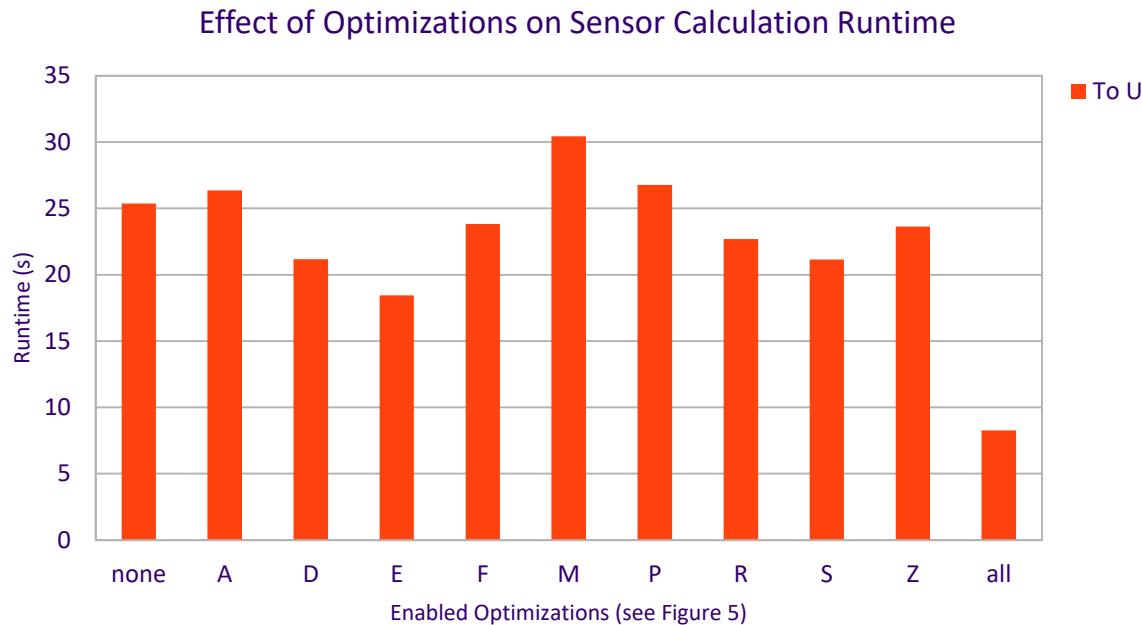
*AfterSum* counts the number of entries in  $C$  after summing colliding partial products together.

# Expressiveness Experiment



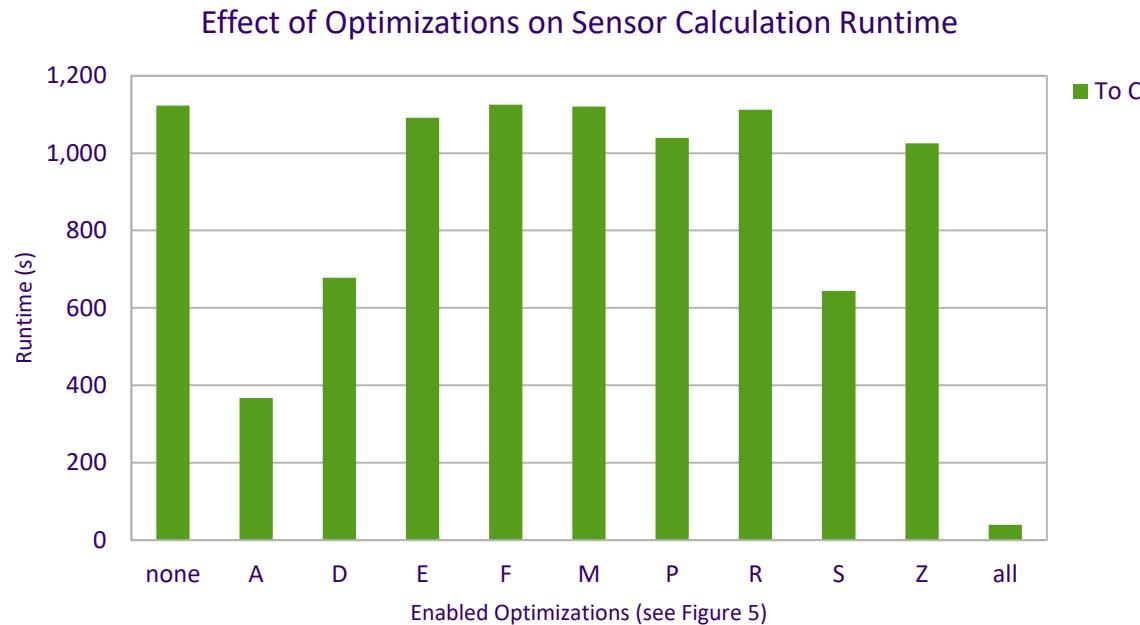
- 4 Worker nodes
- 1.5 months' sensor data
  - Partitioned in 3 day segments
  - 1.2 GB raw; 60 MB after filter to 1 month, project, compress in HDFS

# Expressiveness Experiment



- 4 Worker nodes
- 1.5 months' sensor data
  - Partitioned in 3 day segments
  - 1.2 GB raw; 60 MB after filter to 1 month, project, compress in HDFS

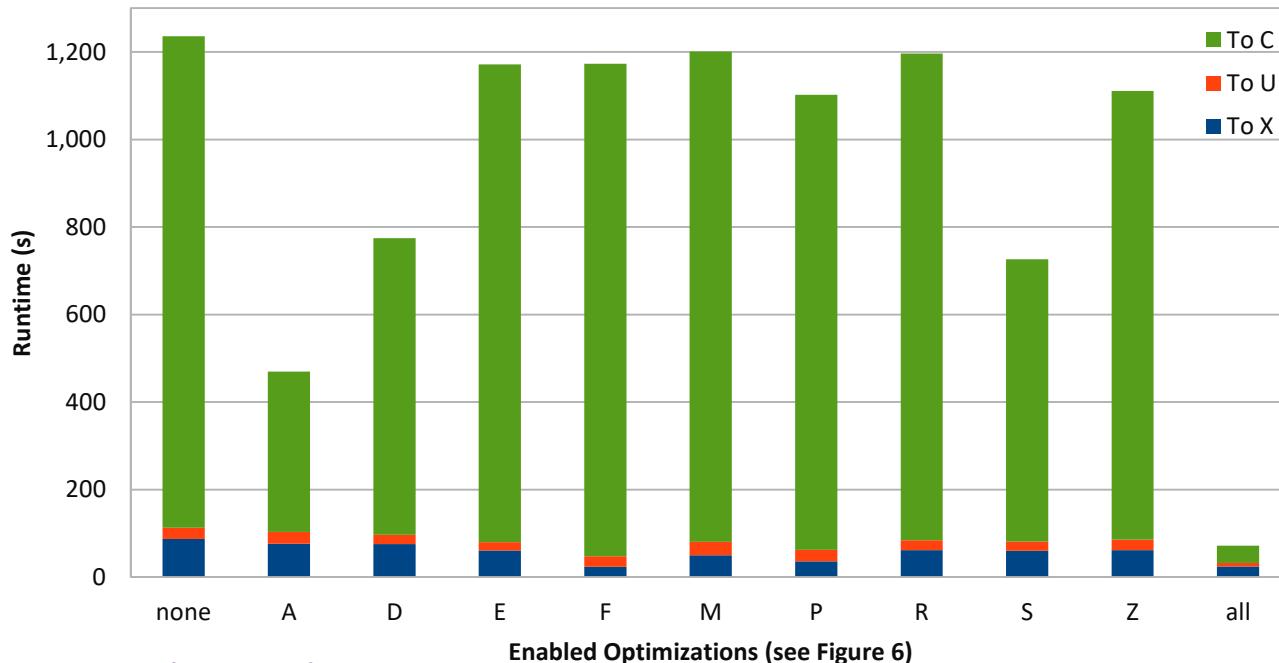
# Expressiveness Experiment



- 4 Worker nodes
- 1.5 months' sensor data
  - Partitioned in 3 day segments
  - 1.2 GB raw; 60 MB after filter to 1 month, project, compress in HDFS

# Expressiveness Experiment

## Effect of Optimizations on Sensor Calculation Runtime



3 Parts of Sensor Calc  
Matrix multiply & divide  
Subtract mean  
Filter, bin, join sensors

- 4 Worker nodes
- 1.5 months' sensor data
  - Partitioned in 3 day segments
  - 1.2 GB raw; 60 MB after filter to 1 month, project, compress in HDFS

Pseudocode	LARA Logical Plan
1 $A, B = \text{LOAD } 's1', 's2'$	$A = \text{LOAD } 's1'$
2 // RA-style (SQL) bin, filter	$A_1 = \text{MAP } A \text{ BY } [v: \text{if}(460 \leq t \leq 860) v \text{ else } \perp]$
3 $A' = \text{SELECT } \text{bin}(t) \text{ AS } t', c, \text{avg}(v)$ FROM $A$ WHERE $460 \leq t \leq 860$	$A_2 = \text{EXT } A_1 \text{ BY } \begin{array}{c cc} t' & v & \text{cnt} \\ \hline \text{bin}(t) & v & v \neq \perp \end{array}$
4 GROUP BY $t', c$	$A_3 = \text{AGG } A_2 \text{ ON } t', c \text{ BY } [v: +, \text{cnt}: +]$
5	$A' = \text{MAP } A_3 \text{ BY } [v: v/\text{cnt}]$
6 $B' = \text{SELECT } \text{bin}(t) \text{ AS } t', c, \text{avg}(v)$ FROM $B$ WHERE $460 \leq t \leq 860$ GROUP BY $t', c$	$B' = \dots // \text{repeat above for second sensor}$
<i>// LA-style (MATLAB) mean, covariance of residuals <math>A' - B'</math>, viewed as dense matrices:</i>	
7 $X = A' - B'$ // residuals; $ t'  \times  c $ matrix	$X = \text{JOIN } A', B' \text{ BY } [v: -]$
8 $N = \text{size}(X, 1)$ // # unique $t'$ s; scalar	$X_1 = \text{MAP } X \text{ BY } [v: v \neq \perp]$
9	$X_2 = \text{AGG } X_1 \text{ ON } t' \text{ BY } [v: \text{any}]$
10	$N = \text{AGG } X_2 \text{ BY } [v: +]$
11 $M = \text{mean}(X, 1)$ //c means; $1 \times  c $ vector	$X_3 = \text{MAP } X \text{ BY } [v: v, \text{cnt}: v \neq \perp]$
12 STORE $M$	$X_4 = \text{AGG } X_3 \text{ ON } c \text{ BY } [v: +, \text{cnt}: +]$
13	$M = \text{MAP } X_4 \text{ BY } [v: v/\text{cnt}]$
14 $U = X - \text{repmat}(M, N, 1)$ //subtract mean	$U = \text{JOIN } X, M \text{ BY } [v: -]$
15 $C = U^T U / (N - 1)$ //c covariances;	$U_1 = \text{RENAME } U \text{ FROM } c \text{ TO } c'$
16 STORE $C$ // $ c  \times  c $ matrix	$U_2 = \text{JOIN } U_1, U \text{ BY } [v: \times]$
17	$U_3 = \text{AGG } U_2 \text{ ON } c, c' \text{ BY } [v: +]$
18	$C = \text{JOIN } U_3, N \text{ BY } [v: v/(v' - 1)]$

$t'$	$c$	$v$	$\perp$	$t'$	$c$	$v$	$\perp$	$c_1$	$c_2$	$v$	$\perp$
460	temp	-3.1	$\perp$	460	temp	0.4	$\perp$	temp	temp	0.32	$\perp$
460	hum	1.6	$\perp$	460	hum	-3.5	$\perp$	temp	hum	0.96	$\perp$
520	temp	-4.0	$\perp$	520	temp	-0.4	$\perp$	hum	temp	0.96	$\perp$
520	hum	-0.8	$\perp$	520	hum	-1.2	$\perp$	hum	hum	2.88	$\perp$

Sensor Query, complete  
translation to LARA

LARA Physical Plan	Access Path	Optimizations
1 $A = \text{LOAD } 's1'$	$[t, c]$	(E) Encode numeric attributes in packed byte form
2 $A_1 = \text{MAP } A \text{ BY } [v: \text{ if}(460 \leq t \leq 860) v \text{ else } \perp]$	$[t, c]$	(F) Push filter into LOAD 's1' FROM 460 TO 860
3 $A_2 = \text{EXT } A_1 \text{ BY } [t': \text{bin}(t) \rightarrow v: v, cnt: v \neq \perp]$	$[t, c, t']$	
3.5 $A_{20} = \text{SORT } A_2 \text{ TO } [t', c, t]$	$[t', c, t]$	(M) Eliminate <b>SORT</b> by $t': \text{bin}(t)$ monotone in $t$
4 $A_3 = \text{MERGEAGG } A_{20} \text{ ON } t', c \text{ BY } [v: +, cnt: +]$	$[t', c]$	
5 $A' = \text{MAP } A_3 \text{ BY } [v: v/cnt]$	$[t', c]$	
6 $B' = \dots // \text{ repeat above for second sensor}$	$[t', c]$	
7 $X = \text{MERGEJOIN } A', B' \text{ BY } [v: -]$	$[t', c]$	(P) Propagate $A, B$ 's partition splits throughout
8 $X_1 = \text{MAP } X \text{ BY } [v: v \neq \perp]$	$[t', c]$	
9 $X_2 = \text{MERGEAGG } X_1 \text{ ON } t' \text{ BY } [v: \text{any}]$	$[t']$	
10 $N = \text{AGG } X_2 \text{ BY } [v: +]$	$[]$	(C) Store scalar $N$ at client instead of a table
10.5 $X_0 = \text{SORT } X \text{ TO } [c, t']$	$[c, t']$	(Z) If $M, C$ relaxed to <i>sparse</i> matrix interpretation, identify $\perp$ with 0, discarding 0-valued entries in $X_3$ and all following tables
11 $X_3 = \text{MAP } X_0 \text{ BY } [v: v, cnt: v \neq \perp]$	$[c, t']$	
12 $X_4 = \text{MERGEAGG } X_3 \text{ ON } c \text{ BY } [v: +, cnt: +]$	$[c]$	
13 $M = \text{MAP } X_4 \text{ BY } [v: v/cnt]$	$[c]$	(D) Defer $X_3, X_4, M$ to future scans on $X_0$ , eliminating write-out of $M$
13.5 $\text{STORE } M$	$[c]$	
14 $U = \text{MERGEJOIN } X_0, M \text{ BY } [v: -]$	$[c, t']$	(R) Reuse $X_0$ data source (common sub-expression)
14.5 $U_0 = \text{SORT } U \text{ TO } [t', c]$	$[t', c]$	$(U_2$ has a similar sub-expression below)
15 $U_1 = \text{RENAME } U_0 \text{ FROM } c \text{ TO } c'$	$[t', c']$	(S) $U^\top U$ is symmetric; only compute upper triangle
16 $U_2 = \text{MERGEJOIN } U_0, U_1 \text{ BY } [v: \times]$	$[t', c, c']$	via MAP filter $c \leq c'$
16.5 $U_{20} = \text{SORT } U_2 \text{ TO } [c, c', t']$	$[c, c', t']$	(A) Push sum of partial products into $U_{20}$ compaction and flush; assume no repeated writes due to server failure
17 $U_3 = \text{MERGEAGG } U_{20} \text{ ON } c, c' \text{ BY } [v: +]$	$[c, c']$	
18 $C = \text{MERGEJOIN } U_3, N \text{ BY } [v: v/(v' - 1)]$	$[c, c']$	(D) Defer $U_3, C$ to future scans on $U_{20}$ , eliminating final pass
18.5 $\text{STORE } C$	$[c, c']$	

### Example optimization:

- ✓ Eliminate **SORT** after monotonic **EXT**

$$(M) \frac{A : k \rightarrow v \quad f(k, v) : k' \rightarrow v' \quad f \text{ monotone in } k}{\text{EXT } A \text{ BY } f \text{ OVER } [k', k]} \quad \text{SORT } (\text{EXT } A \text{ BY } f) \text{ TO } [k', k]$$

# LARADB



t	c	t'	[T] v	[0] cnt
466	temp	460	55.2	1
466	hum	460	40.1	1
492	temp	520	56.3	1
492	hum	520	35.0	1
528	temp	520	56.5	1

*Join*



*Union*



*Extension*



slide from Andreas Kunft, Alexander Alexandrov, Asterios Katsifodimos, Volker Markl's  
2016 BeyondMR talk: Towards optimization across linear and relational algebra

## Introduction

Relational Algebra (RA): Preprocessing  
Linear Algebra (LA): Analytics

Dataflow programs evolved from simple jobs to **complex pipelines**

> These pipelines are composed of **different domains**, e.g.  
collection processing, matrix processing

> Two implementation choices

Some tasks and *properties* easier  
to express in RA, others in LA

Single DSL > **Impedance Mismatch** (SQL OR Matlab)

Multiple DSLs > **Lost optimization potential** (SQL AND Matlab)

Using both RA and LA is difficult without a means of translation,  
such as a *common algebra*

# Last year's Sensor example: Combine measurements, compute means & covariances

```
1 // Read measurements into the DataBags A and B
2 val A = readCSV(...) //
3 val B = readCSV(...)
4 // SELECT a1, ..., aN, b1, ..., bM
5 // FROM A, B
6 // WHERE A.id = B.id
7 val X = for {
8     a <- A
9     b <- B
10    if a.id == b.id
11    } yield (a1, ..., aN, b1, ..., bM)
12
13 // Convert DataBag X into Matrix M
14 val M = X.toMatrix()
15 // Calculate the mean of each column c of the matrix
16 val means = for (c <- M.cols()) yield (mean(c))
17 // Compute the deviation of each cell of M
18 // to the cell's column mean.
19 val U = M - Matrix.fill(M numRows, M numCols)
20             ((i, j) => means(j))
21 // Compute the covariance matrix
22 val C = 1 / (U numRows - 1) * U.t %*% U
```

Preprocessing  
(easier to  
express in RA)

Analysis  
(easier to  
express in LA)

# Use Lara underneath Weld?

---

- > Weld replaced the implementation of Pandas (RAish) and NumPy (Laish)
- > Weld chose a MapReduce algebra
- > Would Weld benefit from using LARA algebra instead?

S. Palkar, J. J. Thomas, A. Shanbhag, D. Narayanan, H. Pirk, M. Schwarzkopf, S. Amarasinghe, M. Zaharia, and S. InfoLab. Weld: A common runtime for high performance data analytics. CIDR, Jan. 2017.



# Sensor Query, Alternate Plan

```
A, B = LOAD('s1.csv'), LOAD('s2.csv')  
A' = SELECT bin(t) AS t', c, avg(v) AS v  
      FROM A WHERE 460 <= t <= 860  
      GROUP BY t', c  
  
B' = ... // repeat above
```

```
X = A - B;  
N = size(X, 1);  
M = mean(X, 1);  
disp(M);  
  
U = X - repmat(M, N, 1);  
C = U.' * U ./ (N - 1);  
disp(C);
```

```
A, B = LOAD('s1.csv'), LOAD('s2.csv')  
A' = SELECT bin(t) AS t', c, avg(v) AS v  
      FROM A WHERE 460 <= t <= 860  
      GROUP BY t', c  
  
B' = ... // repeat above
```

```
X = A - B;  
N = size(X, 1);  
M = mean(X, 1);  
disp(M);  
  
C = N ./ (N - 1) .* ((X.' * X ./ N) - M.' * M);  
disp(C);
```

# Filter as LARA Map (Ext)

t	c	[ $\perp$ ] v
440	temp	54.0
440	hum	40.8
466	temp	55.2
466	hum	40.1
492	temp	56.3
492	hum	35.0
528	temp	56.5



$A' = \text{SELECT } t, c, v \text{ FROM } A$   
 $\text{WHERE } 460 \leq t \leq 860$

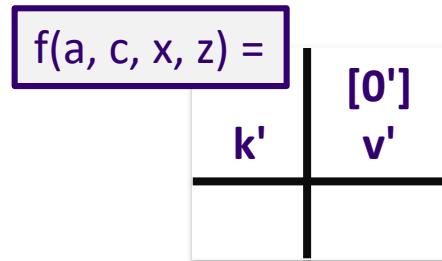
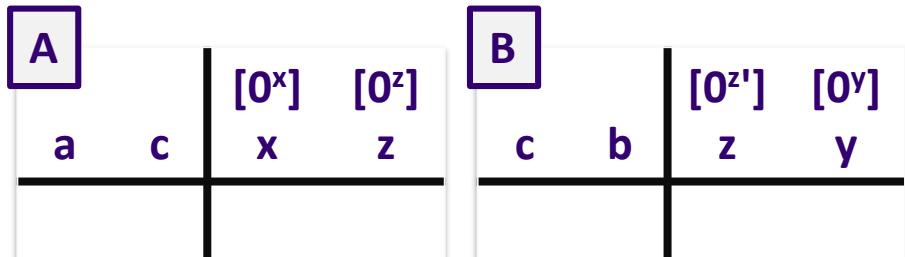
$A_1 = \text{MAP } A \text{ BY}$   
 $[v: \text{if } (460 \leq t \leq 860) v \text{ else } \perp]$

( $)$	$\perp$
( $)$	$\text{if } (460 \leq t \leq 860) v \text{ else } \perp$

t	c	[ $\perp$ ] v
440	temp	$\perp$
440	hum	$\perp$
466	temp	55.2
466	hum	40.1
492	temp	56.3
492	hum	35.0
528	temp	56.5

UDF  $f$  as a table-valued function, ready for ext (flatmap)

# LARA on One Slide



UDFs:  $\oplus : i \times i \rightarrow i$  for  $i \in \{x, y, z\}$

$\otimes : z \times z \rightarrow z'$

$f : a, c \times x, z \rightarrow (k' \rightarrow v' : 0')$

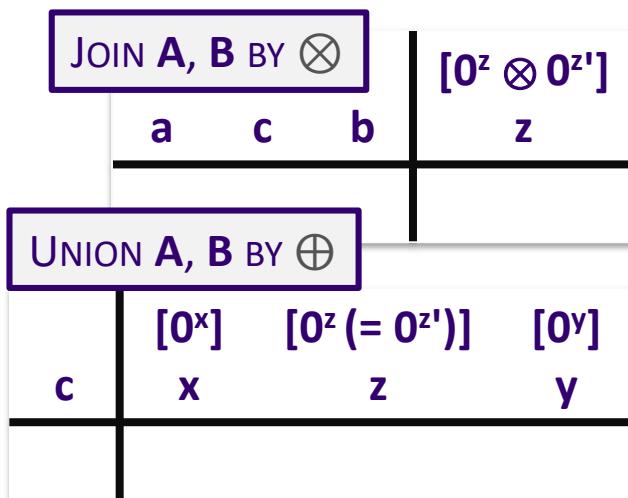
UDF Requirements:

$$\forall i, i \oplus 0^i = 0^i \oplus i = 0^i \text{ for } i \in \{x, y, z\}$$

$$\forall z, 0^z \otimes z = z \otimes 0^z = 0^z \otimes 0^z$$

$$\forall a, c, k', f(a, c, 0^x, 0^z)(k') = 0'$$

$\forall a, c, x, z, f(a, c, x, z)$  has finite support



$A \bowtie_{\otimes} B : a, c, b \rightarrow z : 0^z \otimes 0^z$

$(A \bowtie_{\otimes} B)(a, c, b) := [z : \pi_z A(a, c) \otimes \pi_z B(c, b)]$

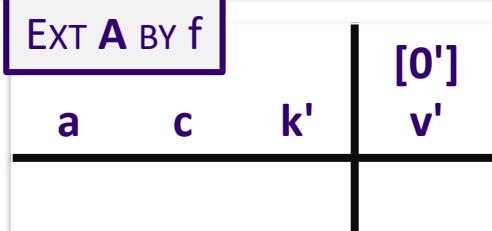
$A \boxtimes_{\oplus} B : c \rightarrow x, z, y : 0^x, 0^z, 0^y$

$(A \boxtimes_{\oplus} B)(c) := [x : \bigoplus_a \pi_x A(a, c),$

$z : \bigoplus_a \pi_z A(a, c) \oplus \bigoplus_b \pi_z B(c, y), y : \bigoplus_b \pi_y B(c, y)]$

$\text{ext}_f A : a, c, k' \rightarrow v' : 0'$

$(\text{ext}_f A)(a, c, k') := f(a, c, A(a, c))(k')$



# Translating LA and RA to LARA

Given  $A : i, j \rightarrow v$ ,  $B : j, k \rightarrow v$ ,  $C : i, k \rightarrow v$

GraphBLAS

		LA	LARA
RA	LARA		
$\sigma_p$	$\text{map}_p$	$A \oplus . \otimes B$	$\boxtimes_{\oplus}^{i,k}(A \bowtie_{\otimes} B)$
$\pi$	map or $\boxtimes$	$A \otimes C$	$A \bowtie_{\otimes} C$
$\times, \bowtie$	$\bowtie$	$A \oplus C$	$A \boxtimes_{\oplus} C$
$\gamma, \cup$	$\boxtimes_{\oplus}$	reduce( $A, \oplus$ )	$\boxtimes_{\oplus}^i A$
(a) RA to LARA		$A(I, J)$	$A \bowtie I \bowtie J$
		$f(A)$	$\text{map}_f$
		$A^\top$	rename

(b) LA to LARA

Roughly: LA, RA  $\subseteq$  LARA  $\subseteq$  MapReduce

# Convert RA Relation, LA Matrix to LARA Associative Table

- 1. Identify keys; separate attributes into keys and values
- 2. Introduce sparsity; make the table a total function: keys → values with finite support

t'	c	v
460	temp	55.2
460	hum	40.1
520	temp	56.4
520	hum	35.0

$$\rightarrow \begin{matrix} & \text{temp} & \text{hum} \\ 460 & [55.2 & 40.1] \\ 520 & [56.4 & 35.0] \end{matrix}$$

t'	c	v
460	temp	55.2
460	hum	40.1
520	temp	56.4
520	hum	35.0

# Convert RA Relation, LA Matrix to LARA Associative Table

1. Identify keys; separate attributes into keys and values
2. Introduce sparsity; make the table a total function: keys → values with finite support

**Support**

$t'$	$c$	$v$
460	temp	55.2
460	hum	40.1
520	temp	56.4
520	hum	35.0
461	temp	⊥
461	wind	⊥
902	hum	⊥
3	wind	⊥
...	...	...

**Set of keys that map to non-default values**

	<i>temp</i>	<i>hum</i>	<i>wind</i>	...
3	⊥	⊥	⊥	...
29	⊥	⊥	⊥	...
460	55.2	40.1	⊥	...
461	⊥	⊥	⊥	...
520	56.4	35.0	⊥	...
557	⊥	⊥	⊥	...
902	⊥	⊥	⊥	...
911	⊥	⊥	⊥	...
...	...	...	...	...

**Default value**

$t'$	$c$	$v$
460	temp	55.2
460	hum	40.1
520	temp	56.4
520	hum	35.0
(3	wind	⊥)
(461	temp	⊥)
(...	...	...)

# Filter as LARA Map (Ext)

t	c	[ $\perp$ ] v
440	temp	54.0
440	hum	40.8
466	temp	55.2
466	hum	40.1
492	temp	56.3
492	hum	35.0
528	temp	56.5

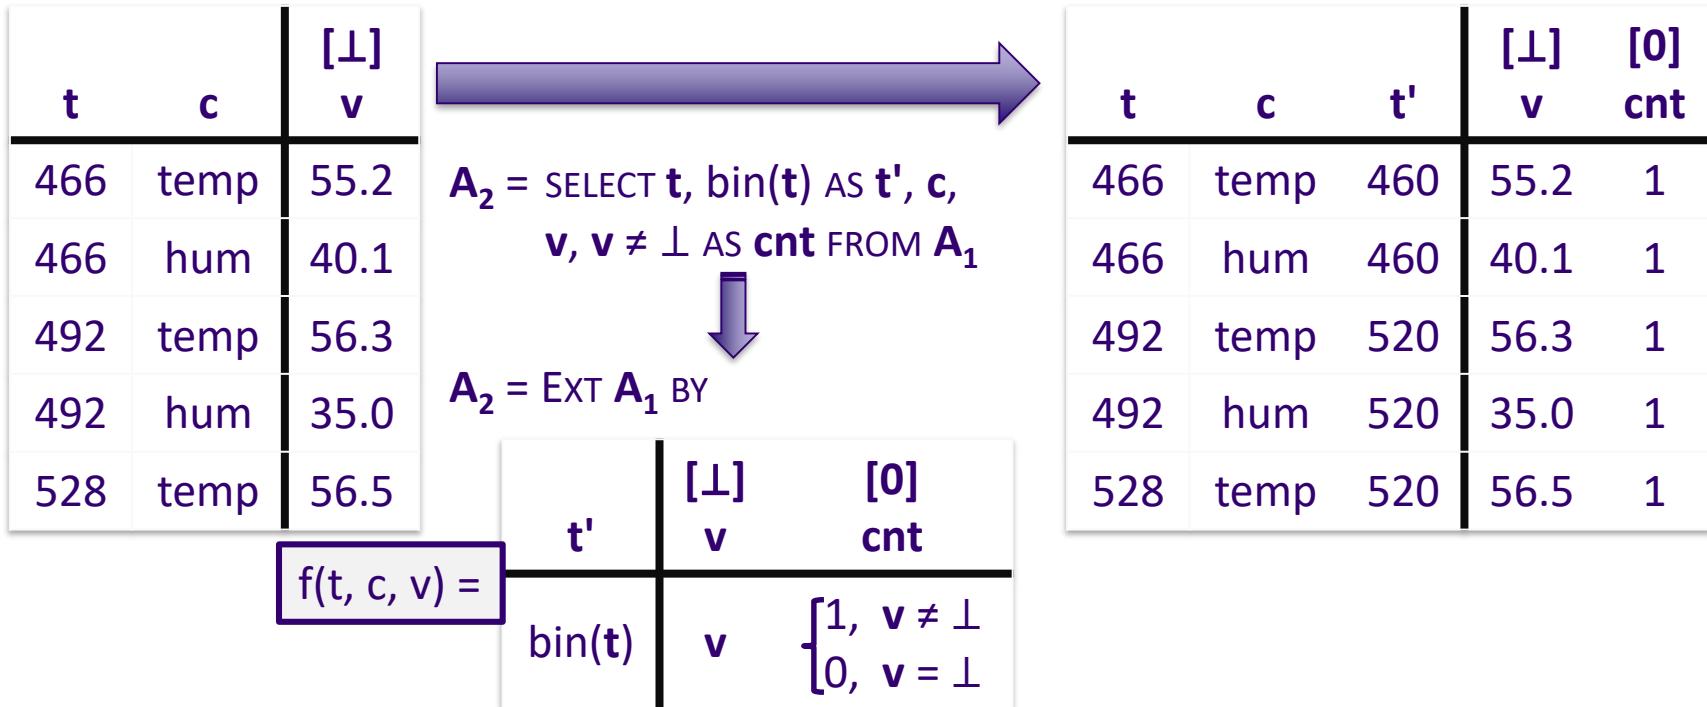


$A' = \text{SELECT } t, c, v \text{ FROM } A$   
 $\text{WHERE } 460 \leq t \leq 860$

$A_1 = \text{MAP } A \text{ BY}$   
 $[v: \text{if } (460 \leq t \leq 860) v \text{ else } \perp]$

t	c	[ $\perp$ ] v
440	temp	$\perp$
440	hum	$\perp$
466	temp	55.2
466	hum	40.1
492	temp	56.3
492	hum	35.0
528	temp	56.5

# Bin as LARA Ext



- Appends new key **t'**, replaces values
- Ext function **f** maintains finite support because
 
$$\forall t, c, t', \quad f(t, c, \perp)(t') = (\perp, 0)$$

# Avg as LARA Agg (Union)

t	c	t'	[⊥] v	[0] cnt
466	temp	460	55.2	1
466	hum	460	40.1	1
492	temp	520	56.3	1
492	hum	520	35.0	1
528	temp	520	56.5	1



$A_3 = \text{SELECT } t', c, \text{sum}(v) \text{ AS } v, \text{sum}(cnt) \text{ AS } cnt \text{ FROM } A_2$



$A_3 = \text{AGG } A_2 \text{ ON } t', c \text{ BY } [v: +, cnt: +]$

t'	c	[⊥] v	[0] cnt
460	temp	55.2	1
460	hum	40.1	1
520	temp	112.8	2
520	hum	35.0	1

- Agg is Union with an empty table →

t'	c	()

- Another Map finishes the calculation of avg(v):  $f(t', c, v, cnt) = [v: v / cnt]$

# Subtract as LARA Join

		[ $\perp$ ]
t'	c	v
460	temp	55.2
460	hum	40.1
520	temp	56.4
520	hum	35.0

$\bowtie_{-}$

		[ $\perp$ ]
t'	c	v
460	temp	58.3
460	hum	38.5
520	temp	60.4

=

		[ $\perp$ ]
t'	c	v
460	temp	-3.1
460	hum	1.6
520	temp	-4.0

$$X = A - B$$



$X = \text{JOIN } A, B \text{ BY } [v: -]$

- > Last year Andreas observed that computational pipelines increasingly cross data models, in particular RA and LA. He illustrated why this is a problem with a sensor example.
- > We found data that made the sensor example real, from a project called the Array of Things started right here in the city of Chicago.
  - Added a few more preprocessing tasks (filtering, binning) to fit the real data.
- > Why is a common algebra possible?
- > Lara; discuss associative table on the spot
- > Join, Union, (Join-Union duality), Ext, properties
- > Sensor example pseudocode, logical Lara
- > Sorted Distributed Maps – relational, key-value, array systems
- > Lower to Physical algebra, Accumulo
- > Admits lots of optimizations; these significantly reduced runtime
- > Does performance take a nosedive if we use Lara as a guiding abstraction for implementations? No; competitiveness experiment with MapReduce, analytics engine natively integrated with Accumulo on MxM, a fundamental building block for LA and RA

# Ext: Flatmap

$$\text{ext}_f$$

$f(a, c, x, z) =$

$k'$	$v'$
$ac$	$x - z$
$ca$	$z - x$

$$= \begin{array}{c|cc|cc} & & [0] & [0] & \\ \hline a & c & x & z & \\ \hline a_1 & c_1 & 11 & 1 & \\ a_1 & c_2 & 12 & 2 & \\ a_2 & c_1 & 13 & 3 & \end{array}$$

$$[0 - 0 = 0]$$

$$v'$$

$$= \begin{array}{c|cc|cc} & a & c & k' & v' \\ \hline a_1 & c_1 & a_1 c_1 & 11 - 5 & \\ a_1 & c_1 & c_1 a_1 & 5 - 11 & \\ a_1 & c_2 & a_1 c_2 & 12 - 2 & \\ a_1 & c_2 & c_2 a_1 & 2 - 12 & \\ a_2 & c_1 & a_2 c_1 & 13 - 3 & \\ a_2 & c_1 & c_1 a_2 & 3 - 13 & \end{array}$$

$$\text{ext}_f A : a, c, k' \rightarrow v' : 0'$$

$$(\text{ext}_f A)(a, c, k') := f(a, c, A(a, c))(k')$$

Requires:

$$\forall a, c, k', f(a, c, 0^x, 0^z)(k') = 0'$$

# Covariance query: easy in LA, harder in RA

$$X = \underbrace{\left[ \begin{array}{cc} x_{11} & x_{12} \\ x_{21} & x_{22} \\ x_{31} & x_{32} \end{array} \right]}_{n \text{ points}} \quad \overbrace{\quad}^{d \text{ attributes}}$$

LA

**N** = size(**X**, 1);

**M** = mean(**X**, 1);

**C** = **X**'\***X** / **N** - **M**'\***M**;

$M = \frac{1}{n} \mathbf{1}^T X$  is a  $1 \times d$  matrix

$C = \frac{1}{n} X^T X - M^T M$  is a  $d \times d$  matrix

RA

(Generated SQL statements for each entry)

**T** = SELECT FROM **X** sum(1.0) AS N,

sum( $X_1$ ) AS  $M_1$ , sum( $X_2$ ) AS  $M_2$ , ..., sum( $X_d$ ) AS  $M_d$ ,

sum( $X_1 * X_1$ ) AS  $Q_{11}$ , sum( $X_1 * X_2$ ) AS  $Q_{12}$ , ...,

sum( $X_{d-1} * X_d$ ) AS  $Q_{(d-1)d}$ , sum( $X_d * X_d$ ) AS  $Q_{dd}$

**C** = SELECT FROM **T**

(1 AS i, 1 AS j,  $Q_{11}/N - M_1 * M_1$  AS v) UNION

(1 AS i, 2 AS j,  $Q_{12}/N - M_1 * M_2$  AS v) UNION

...

# Covariance query: easy in LA, harder in RA

$$X = \underbrace{\left[ \begin{array}{cc} x_{11} & x_{12} \\ x_{21} & x_{22} \\ x_{31} & x_{32} \end{array} \right]}_{n \text{ points}} \quad \overbrace{\quad \quad \quad}^{d \text{ attributes}}$$

LA

**N** = size(**X**, 1);

**M** = mean(**X**, 1);

**C** = **X**'\***X** / **N** - **M**'\***M**;

Mean

$$M = \frac{1}{n} \mathbf{1}^T X$$

Covariance

$$C = \frac{1}{n} X^T X - M^T M$$

RA

(Generated SQL statements for each entry)

**T** = SELECT FROM **X** sum(1.0) AS N,

sum(**X**<sub>1</sub>) AS M<sub>1</sub>, sum(**X**<sub>2</sub>) AS M<sub>2</sub>, ..., sum(**X**<sub>d</sub>) AS M<sub>d</sub>,

sum(**X**<sub>1</sub>\***X**<sub>1</sub>) AS Q<sub>11</sub>, sum(**X**<sub>1</sub>\***X**<sub>2</sub>) AS Q<sub>12</sub>, ...,

sum(**X**<sub>d-1</sub>\***X**<sub>d</sub>) AS Q<sub>(d-1)d</sub>, sum(**X**<sub>d</sub>\***X**<sub>d</sub>) AS Q<sub>dd</sub>

**C** = SELECT FROM **T**

(1 AS i, 1 AS j, Q<sub>11</sub>/N - M<sub>1</sub>\*M<sub>1</sub> AS v) UNION

(1 AS i, 2 AS j, Q<sub>12</sub>/N - M<sub>1</sub>\*M<sub>2</sub> AS v) UNION

...

# LARA Properties

- > If  $\oplus$  or  $\otimes$  is associative, commutative, or idempotent, then so is Union or Join
- > (Push Aggregation into Join) If  $\otimes$  distributes over  $\oplus$ , then

$$(A \bowtie_{\otimes} B) \boxtimes_{\oplus} C = ((\boxtimes_{\oplus}^{k_B \cup k_C} A) \bowtie_{\otimes} (\boxtimes_{\oplus}^{k_A \cup k_C} B)) \boxtimes_{\oplus} (\boxtimes_{\oplus}^{k_A \cup k_B} C)$$

- > (Distribute Join over Union)  
If  $(k_B \Delta k_C) \cap k_A = \emptyset$ , then

$$\begin{aligned} A \bowtie_{\otimes} (B \boxtimes_{\oplus} C) &= \\ (A \bowtie_{\otimes} B) \boxtimes_{\oplus} (A \bowtie_{\otimes} C) & \end{aligned}$$

$$\begin{aligned} & \text{tr}(ABC) \\ &= \boxtimes_+ \text{ext}_{i=l}(ABC) && \text{tr defn.} \\ &= \boxtimes_+ \text{ext}_{i=l} \boxtimes_+^{i,l} (\boxtimes_+^{i,k} (A_{ij} \bowtie_{\otimes} B_{jk}) \bowtie_{\otimes} C_{kl}) && ABC \text{ defn.} \\ &= \boxtimes_+ \boxtimes_+^{i,l} \text{ext}_{i=l} (\boxtimes_+^{i,k} (A_{ij} \bowtie_{\otimes} B_{jk}) \bowtie_{\otimes} C_{kl}) && \text{push ext into } \boxtimes \\ &= \boxtimes_+ \text{ext}_{i=l} (\boxtimes_+^{i,k} (A_{ij} \bowtie_{\otimes} B_{jk}) \bowtie_{\otimes} C_{kl}) && \text{combine } \boxtimes \\ &= \boxtimes_+ (\boxtimes_+^{i,k} (A_{ij} \bowtie_{\otimes} B_{jk}) \bowtie_{\otimes} C_{ki}) && \text{apply ext} \\ &= \boxtimes_+ \boxtimes_+^{i,k} (A_{ij} \bowtie_{\otimes} B_{jk} \bowtie_{\otimes} C_{ki}) && \text{distr. } \bowtie \text{ into } \boxtimes \\ &= \boxtimes_+ (A_{ij} \bowtie_{\otimes} B_{jk} \bowtie_{\otimes} C_{ki}) && \text{combine } \boxtimes \\ &= \boxtimes_+ (B_{jk} \bowtie_{\otimes} C_{ki} \bowtie_{\otimes} A_{ij}) && \text{commute } \bowtie_{\otimes} \\ &= \dots // \text{reversing the above steps} \\ &= \text{tr}(BCA) \end{aligned}$$

= sum(AB  $\otimes$  C<sup>T</sup>)

# Is a common algebra possible?

---

Evidence:

- > Overlapping expressiveness
  - Matrix Multiply = GroupBy  $\circ$  Apply  $\circ$  Join
$$A_{i,j} \oplus.\otimes B_{j,k} = \gamma_{\oplus(v)}^{i,k} \pi_{i,j,k,v=v_A \otimes v_B} (A_{i,j,v_A} \bowtie B_{j,k,v_B})$$
- > Emergence of hybrid systems
  - Spark, Myria, MADlib, SystemML, Weld, ...
  - Systems with RA and LA interfaces and/or data models



# Roadmap

---

- > Motivation: Sensor query
  - Sometimes RA is better, sometimes LA
- > LARA Algebra
  - 3 Operators & Properties
- > LARADB on *Sorted Distributed Maps*
  - Physical operators & optimizations
- > Experiments
  - Expressiveness: LARA affords impactful optimizations
  - Competitiveness: LARA affords competitive performance